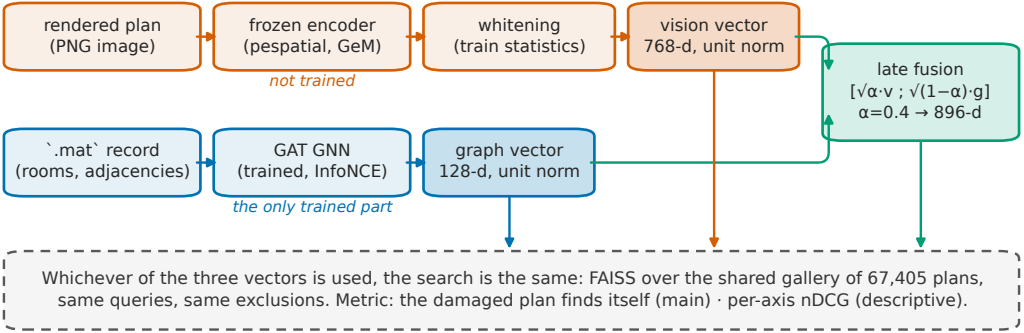
rendered plan
(PNG image)
frozen encoder
(pespatial, GeM)
not trained
whitening
(train statistics)
vision vector
768-d, unit norm
late fusion
[√α·v ; √(1−α)·g]
α=0.4 → 896-d
`.mat` record
(rooms, adjacencies)
GAT GNN
(trained, InfoNCE)
the only trained part
graph vector
128-d, unit norm
Whichever of the three vectors is used, the search is the same: FAISS over the shared gallery of 67,405 plans,
same queries, same exclusions. Metric: the damaged plan finds itself (main) · per-axis nDCG (descriptive).